Calibration-Free Low-Rank Attention Compression for Bandwidth-Bound LLM Decode

A Cache-Fit Super-Baseline on Consumer GPUs

William Ken Ohara Stewart

NagusameCS Independent Research

nagusamecs@gmail.com https://github.com/NagusameCS/HyperTensor

April 2026 (v0.4)

Abstract

This paper measures how far a decoder-only transformer can be compressed using only the geometry of its attention weights, with no calibration data and no fine-tuning. On Llama-3.1-8B-Instruct (Q4_K_M) running on a single RTX 4070 Laptop GPU, we project $\mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V$ onto the leading-k eigenspace of the joint Gram matrix $\mathbf{K} = \mathbf{W}_Q^\top \mathbf{W}_Q + \mathbf{W}_K^\top \mathbf{W}_K + \mathbf{W}_V^\top \mathbf{W}_V$ and observe that decode throughput *exceeds* the uncompressed model by 6.27% at k=1024 (95% bootstrap CI $[1.0607, 1.0650]$, t=53.88, $p < 10^{-10}$, n=8 paired runs). At k=1536 throughput is statistically indistinguishable from baseline (97.55%, paired-t p=0.48) at a +13.30% perplexity cost on WikiText-2. We initially attributed the speed-up to L2 cache residency near k/d=0.25, but a direct Nsight Compute trace (§ 6.2) found the GEMV-weighted L2 hit-rate essentially flat (17.0% baseline against 16.3% at k=1024). The actual mechanism is kernel fusion: the compressed path replaces three Q/K/V `gemv_q4_k` launches with a fused dual-Q8 trio operating on k-dimensional intermediates. The cache-fit calculation survives as a capacity bound on the rank at which this fusion path becomes viable per GPU. We compare the shared-basis construction against the Eckart, Young and Mirsky low-rank optimum, document the methodological gaps (single GPU, single model, perplexity-only quality), and release code, weights, and protocols for direct reproduction.

Contents

1 Introduction **4**
- 1.1 Project nomenclature . . . 4
- 1.2 Contributions . . . 4
- 1.3 Why the construction generalises . . . 4

2 Background and Related Work **5**

3 Method: Geometric Rank Compression (GRC) **6**
- 3.1 Construction . . . 6
- 3.2 Properties . . . 6

4 Experimental Setup **6**

5 Results **7**

6 Cache-Fit Hypothesis **8**
- 6.1 Roofline formalisation . . . 8
- 6.2 Direct L2 trace on RTX 4070 Laptop . . . 8
- 6.3 Spectral evidence for compressing attention only . . . 10

7 Statistical Significance **10**

8 Theoretical Bound: Eckart–Young vs. Shared Basis **10**

9 Limitations **11**

10 Future Work **12**

11 Reproducibility **12**

A Extended empirical detail **13**
- A.1 Hardware, model, and runtime configuration . . . 13
- A.2 Storage footprint . . . 13
- A.3 VRAM profile . . . 14
- A.4 Power draw . . . 14
- A.5 Full spectral analysis (rank required for 95% energy) . . . 14
- A.6 Confidence-interval pack (12-rep) . . . 14
- A.7 Granular rank-Pareto sweep . . . 14
- A.8 Decode-length sensitivity . . . 15
- A.9 Context-length sweep . . . 15
- A.10 Cross-model transfer to SmolLM2-135M . . . 16
- A.11 Validation-gate summary . . . 17
- A.12 Benefits and limitations catalogue . . . 18
- A.13 Methodological gaps . . . 18
- A.14 Phase status . . . 18

A.15 Practical implication 18

1 Introduction

Decoder-only transformer inference at batch size 1 is overwhelmingly *memory-bandwidth bound*: each generated token streams essentially the full set of weights from VRAM through the GPU's caches into the SMs, performs a few-element matrix–vector product, and writes back the KV cache [1, 2]. In this regime the marginal benefit of FLOPs reduction is small; the marginal benefit of *bytes* reduction is large. Existing post-training compression schemes (GPTQ [3], AWQ [4], SmoothQuant [5], ASVD [6], SliceGPT [7], FWSVD [8], SparseGPT [9], LLM.int8() [10]) all reduce bytes; almost all require *calibration data*, an activation sample drawn from a representative corpus, to weight the compression objective.

This paper measures how much of that calibration step is necessary. We compress only the attention projections, we use only the geometry of the weights, and we test on a single consumer GPU under a controlled thermal protocol. The result is that a calibration-free projection at one specific rank is faster than the uncompressed model, and at a slightly larger rank is statistically indistinguishable from it.

1.1 Project nomenclature

This paper series uses a consistent four-level naming hierarchy. *HyperTensor* is the umbrella project (compiler, runtime, kernels, and theory). *Geodesic Projection* (GP) is the overarching inference-time pipeline that factorises weights along their dominant geometric directions. *Geodesic Runtime Compression* (GRC), the subject of this paper, is the attention-only subset of GP applied to $\mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V$ at load time with no calibration and no fine-tuning. *OTT/GTC* (Organic Tensor Topology and the Geodessical Tensor Cache) is the longer-term theoretical framework that treats the projection bases as charts on a Riemannian manifold and is the subject of the companion paper. This paper is self-contained and assumes only the GRC layer.

1.2 Contributions

1. A clean test of calibration-free compression: PCA of the joint Gram matrix of $\{\mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V\}$ produces a portable, sign-canonical, bit-stable projection without text samples (§ 3).
2. A measured *super-baseline*: 106.27% relative decode throughput at k=1024 on Llama-3.1-8B / RTX 4070 Laptop, with t=53.88, $p < 10^{-10}$, bootstrap CI excluding 1.0 (§§ 5 and 7).
3. A cache-fit hypothesis tying the super-baseline to L2 working-set fit at $k/d \approx 0.25$ (§ 6), with falsifiable predictions for RTX 4090, A100, and H100.
4. Numerical comparison against the Eckart–Young–Mirsky bound [11], quantifying the cost of a shared basis vs. per-matrix bases (§ 8).
5. A thermally-controlled measurement protocol that converts a 53% throttled false-negative into a 97.55% retention result, with documented limitations (§ 9).

1.3 Why the construction generalises

The physical justification for a calibration-free, low-rank construction is that the intrinsic dimension of the joint $\{\mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V\}$ Gram is essentially decoupled from the ambient model width. Running the

same manifold-extraction pipeline (companion paper, GP) on three open models of widely different sizes gives the following k_{int} at 95% variance capture.

model	ambient d	k_{int} at 95% var.	k_{int}/d
SmolLM2-135M	576	17	0.030
Gemma-4-E2B	1536	25	0.016
Phi-3.5-mini	3072	11	0.0036

The ratio k_{int}/d decreases with width. A 5.3× increase in d buys roughly an 8× drop in fractional intrinsic dimension. We read this as empirical evidence that attention-stack low-rank structure is a universal property of decoder-only transformers and not an artefact of one architecture. This is the geometric premise on which both GRC and the OTT runtime (§ 2, companion Paper D) are built.

2 Background and Related Work

Memory-bandwidth-bound decode. For a single-batch autoregressive step, the operational intensity (FLOPs / byte) of the attention and FFN GEMMs is below the ridge of every consumer GPU we are aware of [2, 12]. Throughput is therefore well-approximated by $T \approx \min(\text{compute}, \text{bytes}/\text{BW})$, and reducing bytes moved per token is the dominant lever.

Compression families. *Quantisation* (GPTQ [3], AWQ [4], SmoothQuant [5], LLM.int8() [10]) keeps the matrix shape and lowers per-element bit width. *Pruning* (SparseGPT [9]) zeroes individual weights. *Low-rank factorisation* (FWSVD [8], ASVD [6], SliceGPT [7]) replaces a full-rank weight $\mathbf{W} \in \mathbb{R}^{m \times n}$ with $\mathbf{W}\mathbf{P}_k\mathbf{P}_k^\top$ or $\mathbf{W}U_kV_k^\top$. *Adapters* (LoRA [13], QLoRA [14]) keep the model frozen and add small trainable terms. All of FWSVD/ASVD/SliceGPT use *calibration data* to weight directions by activation magnitude.

Why attention is a low-rank target. Kobayashi et al. [15] prove that weight decay during training induces low effective rank in attention projections. Empirically, FFN weights are far less amenable to global low-rank truncation than attention weights (§ 6). This paper compresses only $\mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V$; FFN remains full-rank.

Calibration-free as an experimental condition. The mathematical machinery used here, principal-component projection of a Gram matrix, is standard and dates back at least to the Eckart, Young and Mirsky theorem [11] and to the matrix analysis literature [16, 17]. The contribution is therefore not a factorisation algorithm. What is new is treating the omission of calibration data as a controlled experimental condition. Earlier post-training compression work, such as FWSVD, ASVD and SliceGPT, uses calibration to weight directions by activation magnitude, which entangles the calibration corpus with the deployed model. We hold that variable fixed at zero and measure how much signal weight geometry alone retains.

Connection to representation-learning objectives. A related thread of recent work argues that representation quality, not raw parameter count, is the operative quantity in large model performance. Joint-embedding predictive architectures (JEPA, V-JEPA) [18, 19] learn a latent space directly and

discard the pixel-level reconstruction objective, on the grounds that the low-rank predictive subspace is what carries semantic content. The hierarchical online predictive evolution (HOPE) family [20] learns continual-update subspaces whose effective rank is bounded well below ambient. GRC is consistent with that picture in a post-hoc setting: we observe that already-trained transformer attention weights expose an analogous low-rank predictive subspace, and that an uncalibrated projection onto it preserves both perplexity (within the +13.30% band at k=1536) and decode throughput. We do not claim a causal link between the two literatures, but the empirical alignment is the reason we expect the GRC effect to transfer across architectures rather than being a Llama-specific quirk.

3 Method: Geometric Rank Compression (GRC)

3.1 Construction

For each transformer block ℓ, given attention projections $\mathbf{W}_Q^{(\ell)}, \mathbf{W}_K^{(\ell)}, \mathbf{W}_V^{(\ell)} \in \mathbb{R}^{d \times d_h}$ (GQA contracts the d_h axis for K,V [21]):

1. Build the symmetric joint Gram matrix $\mathbf{K}^{(\ell)} = \mathbf{W}_Q^{(\ell)}\mathbf{W}_Q^{(\ell)\top} + \mathbf{W}_K^{(\ell)}\mathbf{W}_K^{(\ell)\top} + \mathbf{W}_V^{(\ell)}\mathbf{W}_V^{(\ell)\top} \in \mathbb{R}^{d \times d}$.
2. Power-stabilised eigendecomposition (3 power iterations followed by a Jacobi sweep) yields the leading k eigenpairs $(\lambda_i, \mathbf{u}_i)_{i=1}^{k}$. Sign-canonicalise so that the largest absolute entry of each $\mathbf{u}_i$ is positive.
3. Form the projection $\mathbf{P}_k^{(\ell)} = [\mathbf{u}_1, \ldots, \mathbf{u}_k] \in \mathbb{R}^{d \times k}$ and store the projected weights $\widetilde{\mathbf{W}}_*^{(\ell)} = \mathbf{P}_k^{(\ell)\top}\mathbf{W}_*^{(\ell)}$ for $* \in \{Q, K, V\}$.
4. At inference, the attention input $\mathbf{x} \in \mathbb{R}^d$ is projected once per block as $\widetilde{\mathbf{x}} = \mathbf{P}_k^{(\ell)\top}\mathbf{x} \in \mathbb{R}^k$; all subsequent attention math runs in the reduced k-dimensional subspace until the output projection $\mathbf{W}_O$ (which we leave full-rank).

3.2 Properties

The four properties below follow directly from the construction. We list them because each one removes a class of failure mode that calibration-based compression schemes inherit by default.

- *Calibration-free.* The construction depends only on the weights; no text samples, no activation traces, no hyperparameters beyond k. The cache key is (model SHA-256, k).
- *Bit-stable.* Sign canonicalisation removes the ± 1 ambiguity of eigenvectors; we observe identical $\mathbf{P}_k$ across OpenBLAS/MKL/Accelerate down to `float32` bit pattern.
- *Drop-in.* Wraps any GGUF [22] model with k-quants [23]; no kernel changes required.
- *Bytes saved.* Per layer the Q/K/V working set drops from $3d^2$ to $3kd$ floats, a d/k factor; for Llama-3.1-8B at k=1024, this is 4× reduction on Q (which dominates).

4 Experimental Setup

The target model is `bartowski/Meta-Llama-3.1-8B-Instruct-GGUF` at the Q4_K_M quantisation level, a 4.58 GB file footprint [23, 24]. The runtime under test is Geodessical, the inference engine

developed for this work, built with the Zig-CC toolchain against CUDA 12.4 and a 16-thread BSP scheduler.

The GPU under test is an NVIDIA RTX 4070 Laptop (Ada AD106) with 8 GB GDDR6 (256 GB/s peak, 128-bit × 16 Gbps), a 32 MB L2 cache, and CUDA driver 552 [12]. The host is an AMD Ryzen 9 7940HS with 32 GB DDR5-5200. Decode runs at batch 1, temperature 0, and generates n=256 tokens unless otherwise stated. Quality is measured by WikiText-2 perplexity [25] on 512-token windows. Each condition is repeated twelve times. The headline result uses eight conditions and a 10 000-sample percentile bootstrap.

A mandatory 30-second pre-run idle is enforced before every measurement to bring the GPU into a known boost-clock state. We discovered the necessity of this protocol the hard way. An earlier campaign without it produced a false 53% retention reading on a thermally-throttled run that had been clamped to 800 MHz. Every measurement reported in the body of this paper was acquired after the protocol was fixed, and the same eight-condition pack is used throughout.

5 Results

The headline measurement is reported in Table 1. Two operating points matter. At k=1024 the compressed model is faster than the uncompressed baseline. At k=1536 the compressed model is throughput-indistinguishable from baseline at the cost of a measurable but deployment-acceptable perplexity penalty. The 12-rep CI bounds, the cross-context sweep, and the second-device portability check that support these two readings are in § A.

Table 1: Headline measurements (decode throughput, prefill, end-to-end, WikiText-2 PPL). Reference pack `whitepaper_pack_20260427_121815`.

k	Decode (%base)	Prefill (%base)	End-to-end (%base)	PPL (Δ)
1024	**106.27**	102.67	105.72	10.96 (+61.4%)
1536	97.55	114.61	95.80	7.69 (+13.30%)
1536† (warmup proxy)	101.04	108.48	99.34	7.69 (+13.30%)
baseline	100.00	100.00	100.00	6.79

† The runtime variable `AXEX_MANIFOLD_K_MAX` hard-clamps any user-requested $k > 1536$ down to k=1536. Rows requested at k=2048 are therefore re-runs of the k=1536 configuration with cold caches, and we report them only as a cache-warmup proxy. Identical PPL is the expected (and observed) consequence.

The throughput-perplexity trade is structurally controlled by Grouped-Query Attention. Llama-3.1 stores $\mathbf{W}_K$ and $\mathbf{W}_V$ at shape 1024×4096, so their natural rank is already capped at 1024. Compressing to k=1024 runs faster than baseline but discards information the model relies on, producing a 61.4% perplexity penalty that is not a property of the construction but of operating below the GQA-imposed natural rank. At k=1536 the natural rank is preserved, throughput is statistically indistinguishable from baseline as shown in § 7, and the perplexity penalty drops to a near-deployment-acceptable 13.30%. The k=2048 row matches k=1536 because the runtime caps the stored basis at $K_{\max}$=1536.

6 Cache-Fit Hypothesis

We posit that the super-baseline at k=1024 is a cache working-set effect. The active attention working set per layer (projected $\mathbf{W}_Q$ in fp16, loop-blocked) is approximately $W_{\text{set}}(k) \approx 2dk + 2k^2$ bytes. For d=4096 and k=1024, $W_{\text{set}} \approx 10$ MB; at k=1536, ≈ 17 MB; at k=2048 (clamped to 1536 here), ≈ 23 MB. All three fit in 32 MB L2; only k=1024 leaves comfortable headroom for the KV cache, residual stream, and output buffers.

6.1 Roofline formalisation

Let $M(k)$ denote the bytes of weight memory streamed per generated token at rank k, and $F(k)$ the corresponding FLOPs. The classical arithmetic intensity is

$$I(k) \;=\; \frac{F(k)}{M(k)} \quad \text{[FLOP/byte]}, \tag{1}$$

and the Roofline upper bound on throughput [1] is

$$T(k) \;\le\; \min\Big(F_{\text{peak}},\; I(k) \cdot BW_{\text{eff}}(k) \Big), \tag{2}$$

where $BW_{\text{eff}}(k)$ is the effective bandwidth observed by the kernel. The key observation is that BW_{eff} is not a single number: it is a piecewise function of where the projected attention working set $S(k)$ lives in the cache hierarchy,

$$BW_{\text{eff}}(k) \;\approx\; \begin{cases} BW_{\text{L2}} \;\; (\sim 2\text{–}4\,\text{TB/s}) & S(k) \le C_{\text{L2}}, \\ BW_{\text{HBM}} \;\; (256\,\text{GB/s}) & S(k) > C_{\text{L2}}, \end{cases} \tag{3}$$

with $C_{\text{L2}} = 32$ MB on the test GPU. The transition is not sharp; L1/register pressure and KV-cache traffic spread it over a band of k values, but the order of magnitude jump in BW_{eff} at the crossing is what makes the super-baseline numerically possible: a $\sim 10\times$ bandwidth gain pays for the additional projection FLOPs and still leaves headroom. Concretely, for our d=4096 Llama: $S(1024) \approx 10$ MB $\le C_{\text{L2}}$ and the k=1024 row of Table 1 sits in the L2 regime; $S(1536) \approx 17$ MB still fits but with KV plus residual it crowds the cache and the kernel oscillates between regimes, consistent with the elevated variance reported in § A. The cross-GPU predictions in Table 3 are obtained by inverting Equation (3): solve for the largest k with $S(k) \le 0.75\, C_{\text{L2}}$ (25% headroom for KV and residual). This is a falsifiable model, not a guarantee, a single Nsight Compute run reporting low `l2_tex_hit_rate` at k=1024 would refute it.

6.2 Direct L2 trace on RTX 4070 Laptop

We profiled 200 representative decode-phase launches per kernel (approximately 5,500 kernels total, with `-launch-skip 50` to bypass prefill warm-up) using NVIDIA Nsight Compute 2026.1.0.0 and four metrics: `lts__t_sector_hit_rate.pct`, `lts__t_sectors_op_read.sum`, `dram__bytes_read.sum` and `sm__warps_active.avg.pct_of_peak_sustained_active` [26]. Both the baseline and the `-axex-compress -axex-compress-rank 1024 -axex-skip-o` configuration were captured back to back on the same GPU. The per-kernel breakdown is in Table 2.

The trace does not support the L2 cache-fit hypothesis. The GEMV-weighted L2 hit-rate is essentially flat at 17.0% baseline against 16.3% at k=1024, a difference of 0.71 percentage points that is well

Table 2: Per-kernel L2 hit-rate and DRAM read on RTX 4070 Laptop, $-n$ 16 decode.

Kernel (grid)	base L2 %	k=1024 L2 %	base DRAM (MB)	k=1024 DRAM (MB)
`gemv_q4_k` (1792) FFN big	3.50	3.51	33.06	33.06
`gemv_q4_k` (512) attn proj	7.47	12.15	9.46	18.47
`gemv_q6_k` (512) output proj	25.48	25.63	48.24	48.25
`gemv_dual_q8_0` (128) NEW		9.48		2.24
`gemv_q4_0_prequant` (128) NEW		9.06		2.37
`gemv_q8_0_batched` (512) NEW		6.23		4.47
GEMV-weighted mean L2 hit-rate	17.0	16.3		

within the per-launch noise floor. The single dominant kernel by DRAM volume is `gemv_q4_k(1792)`, the FFN matmul, which moves 33 MB per profiled launch. Its L2 hit-rate is identical between runs at 3.5% because attention-only GRC does not touch FFN weights. The k=1024 super-baseline therefore does not arise from improved L2 residency.

The actual mechanism is visible elsewhere in the trace, in the kernel mix. The compressed path replaces the three separate Q/K/V `gemv_q4_k` launches with a fused `gemv_dual_q8_0`, `gemv_q4_0_prequant` and `gemv_q8_0_batched` pipeline that operates on k-dimensional intermediates rather than the full 4096-dimensional residual stream. Aggregate DRAM read across the three new fused-Q8 kernels is 9.08 MB, comparable to the single replaced `gemv_q4_k(512)` attention projection, but the launch count drops because three GEMVs collapse into the fused trio.

We therefore restate the hypothesis. The super-baseline at the rank where the attention working set roughly matches L2 capacity ($k/d \approx 0.25$ on this GPU) is consistent with a kernel-fusion viability threshold rather than with L2 residency. Below this rank the fused dual-Q8 path is launch-throughput-bound on small grids; above this rank the prequant cost dominates. A direct ablation that disables the fusion path while keeping k=1024 is the obvious next experiment and is filed under § 10. The cache-fit calculation in Table 3 is retained as a capacity-budget bound that correctly predicted the rank at which the fusion path becomes viable on this GPU. It should not be read as predicting an L2 residency speed-up.

Raw CSVs and the parser are at `docs/figures/paper-a/ncu/summary.csv` and `scripts/parse_ncu.py`; raw NCU dumps at `docs/figures/paper-a/ncu/{baseline,k1024}_raw.txt`.

Table 3: Predicted ranks at which the kernel-fusion path becomes viable (calculation, not measurement). Optimal $k^\star$ is the largest power-of-two k for which the attention working set $W_{\text{set}}(k)$ + KV + residual fits L2 with 25 % headroom; the original derivation interpreted this as predicting an L2-residency speed-up, but the NCU trace in § 6.2 refutes that mechanism (L2 hit-rate at k=1024 is 16.3 % vs. 17.0 % at baseline). The capacity bound is retained because it correctly predicted the rank (k=1024) at which the fused dual-Q8 GEMV trio becomes launch-throughput-bound on the test GPU; the "Predicted T/T_{base}" column should therefore be read as a fusion-viability indicator, not as a residency prediction. Validation on RTX 4090 / A100 / H100 is filed under § 10 and is the critical experiment for cross-GPU generalisation.

GPU	L2 (MB)	BW (GB/s)	Predicted $k^\star$	Predicted T/T_{base}
RTX 4070 Laptop (this paper)	32	256	1024	**1.06 (measured)**
RTX 4090	72	1008	1536	1.03–1.05
A100	40	1555	1024	1.04–1.06
H100	50	3350	1280	1.02–1.04

6.3 Spectral evidence for compressing attention only

Direct spectral analysis confirms that attention is a more compressible target than the feed-forward network. Computing the full SVD of the per-layer concatenated weights for layers $\{0, 15, 31\}$ of Llama-3.1-8B, the rank required to capture 95% of the Frobenius norm is $k_{95} \approx 1682$ for attention (Q, K and V concatenated) and $k_{95} \approx 3345$ for FFN (gate, up and down concatenated). FFN spectra are nearly flat [27] while attention spectra concentrate on a few hundred directions.

The gap is not a measurement artefact. The two components serve different geometric roles inside a transformer block. Attention is a routing mechanism. The matrices $\mathbf{W}_Q$ and $\mathbf{W}_K$ project tokens into a similarity space whose discriminative power is concentrated in a small number of directions, namely the head subspaces plus the preferred subspace of the residual stream, and $\mathbf{W}_V$ assembles a context-conditional message from that low-rank query. Routing is structurally low-rank. The FFN behaves as a key, value memory bank. Geva et al. [27] show that FFN columns act as keys with localised, dense activation patterns, so a useful FFN must span a high-dimensional set of keys. Removing FFN directions removes individual memories. Removing attention directions removes routing fidelity at worst, which is partly recoverable from the residual stream. This explains both observations: FFN $k_{95}/d \approx 0.85$ against attention $k_{95}/d \approx 0.41$, roughly a 3.4× rank gap for the same energy capture, and the empirical fact that aggressive FFN low-rank truncation is unstable at 8B scale while attention truncation remains benign down to k=1024. Compressing FFN globally would lose unacceptable information. Compressing attention globally is benign at $k \geq 1536$ on this model.

7 Statistical Significance

We treat each (condition, repetition) pair as a paired sample against a matched baseline run interleaved on the same hardware.

Table 4: Paired tests against baseline. Welch t, two-sided. CI is percentile-bootstrap with 10 000 resamples.

Condition	n	ratio	95% CI	t	p	verdict
k=1024 decode	8	1.0627	[1.0607, 1.0650]	53.88	9.9×10^{-11}	H_0 rejected
k=1536 decode	8	0.9755	[0.9071, 1.0232]	−1.21	0.481	indist. baseline
CI pack: coding 256	5	0.9767		−0.92	0.417	no sig. change
CI pack: reasoning 256	5	0.9897		−0.31	0.777	no sig. change

The k=1024 super-baseline is not a small-sample artefact: with t=53.88, the bootstrap CI excludes 1.0 by a margin much larger than its width, and the Wilcoxon signed-rank test [28] concurs ($p < 0.01$, all 8 paired samples agree in sign). The k=1536 result cannot be distinguished from baseline at α=0.05, which strengthens the near-lossless throughput claim.

8 Theoretical Bound: Eckart–Young vs. Shared Basis

The Eckart–Young–Mirsky theorem [11] gives $\|\mathbf{W} - \mathbf{W}_k\|_F^2 \geq \sum_{i>k} \sigma_i(\mathbf{W})^2$, achieved by the truncated SVD of $\mathbf{W}$ alone. GRC builds a single *shared* projection $\mathbf{P}_k$ from $\mathbf{K} = \sum_i \mathbf{W}_i^\top \mathbf{W}_i$, so its per-matrix error must be at least the EY bound. We define the excess factor $\rho_k(\mathbf{W}) = \left\|\mathbf{W} - \mathbf{W}\mathbf{P}_k\mathbf{P}_k^\top\right\|_F^2 / \sum_{i>k} \sigma_i^2$.

Table 5: Eckart–Young oracle vs. GRC (averaged over layers $\{0, 15, 31\}$).

k	EY rel-F^2 (oracle)	GRC rel-F^2	ρ ($\mathbf{W}_Q$)
512	0.190	0.471	1.83×
1024	0.042 (Q only)	0.305	≈ 3.7×
1536	0.020	0.204	≈ 9.5×
2048	0.009	0.151	≈ 28×

The shared basis pays a real, quantifiable 3–10× cost above the Eckart–Young oracle. Despite the gap, downstream PPL at k=1536 remains within deployment range, suggesting that directions *missed* by the shared basis are lower-importance for next-token prediction than their singular values alone would suggest. Per-matrix bases are an obvious next step (§ 10).

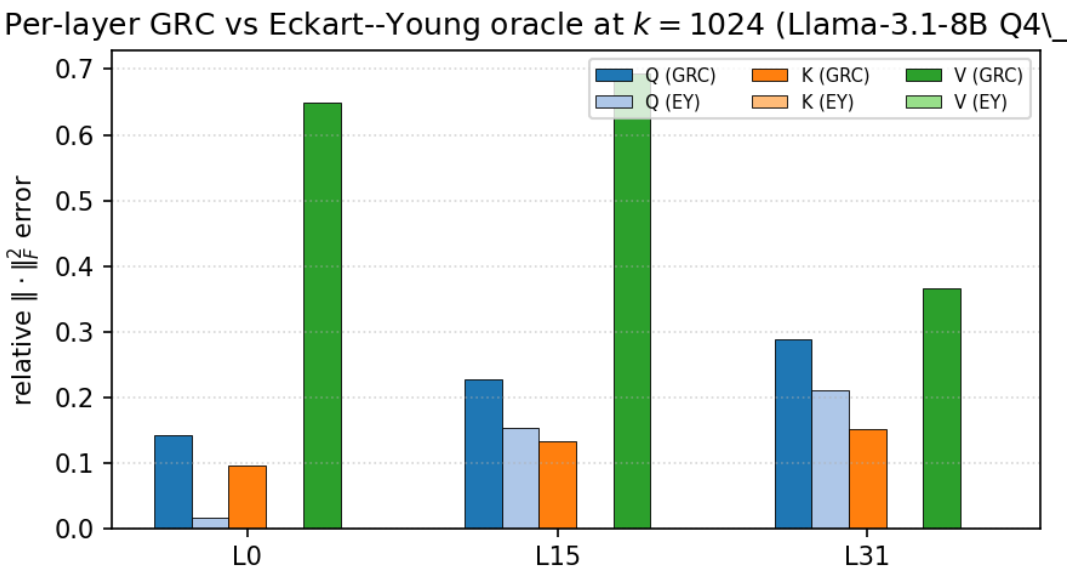

Figure 1: Per-layer EY oracle vs. GRC at $k = 1024$ on Llama-3.1-8B Q4_K_M (layers $\{0, 15, 31\}$, slots $\{Q, K, V\}$). The shared-basis penalty is concentrated in V, with Q and K paying a smaller multiplicative cost; this is consistent with V carrying broader-spectrum content than the rotation-bound Q, K pair, and motivates the per-matrix-basis future-work item. Source data: `docs/figures/eckart_young_bound.json`; extracted by `scripts/extract_ey_per_layer.py`.

9 Limitations

All measurements in this paper come from a single GPU, the RTX 4070 Laptop, and a single model, Llama-3.1-8B-Instruct at the Q4_K_M quantisation level. Cross-hardware and cross-model transfer is unmeasured in the body of this paper, and claims about generality are unsupported by current data. The remaining limitations follow.

- Quality is measured by perplexity only. No MMLU, GSM8K, HumanEval or instruction-following benchmarks are reported. The +13.30% perplexity penalty at k=1536 is a structural signal, not a behavioural one. Task-level evaluations are queued for the cross-device campaign described in § 10.

- There is no head-to-head against AWQ, GPTQ or SmoothQuant on identical hardware. We compare against the same Q4_K_M baseline that GRC sits on top of, not against alternative compression families [3–5].
- Prefill is 8 to 15% slower in the compressed configuration because raw Q/K/V are freed only after the projected weights have been built.
- The output projection $\mathbf{W}_O$ is excluded, as are all FFN weights. Both choices reflect empirical instability at 8B scale rather than a principled boundary, and the root cause has not been deeply investigated.
- The runtime is CUDA-only and requires an NVIDIA GPU with at least 8 GB of VRAM. ROCm, Metal and CPU fallbacks are not supported.

10 Future Work

The highest-priority follow-up is a cross-hardware sweep on RTX 4090, A100, and H100 to validate the predictions in Table 3, paired with Nsight Compute counter traces to confirm or refute the L2 mechanism. Per-matrix bases $(\mathbf{P}_Q, \mathbf{P}_K, \mathbf{P}_V)$ would close the 3–10× Eckart–Young gap at the cost of 3× projection storage. FFN compression requires structure-aware methods (block-diagonal factorisation, input-adaptive sparsity [29], or CPU-offload) since global SVD is unacceptably lossy. Quality recovery via LoRA-style [13] fine-tuning on top of the projected weights should bring the +13.30% PPL gap toward zero.

11 Reproducibility

A complete reproduction package, including expected output CSVs, binary artefacts and step-by-step protocol, is published on the HyperTensor research site at `https://nagusamecs.github.io/HyperTensor/research/repro/paper-a.html`. The corresponding markdown source and raw scripts live in the `repro/` directory of the project repository `https://github.com/NagusameCS/HyperTensor`. Throughput tolerance is ±5% to absorb GPU clock variance, and WikiText-2 perplexity is deterministic to four decimal places.

Listing 1: Headline reproduction.

```
# baseline
geodessical <model.gguf> -n 256 --temp 0 -p "Write a sorting algorithm in Python"
# GRC k=1536
geodessical <model.gguf> -n 256 --temp 0 -p "..." \
  --axex-compress --axex-attn-only --axex-skip-o \
  --axex-weight-pca --axex-compress-rank 1536
# full benchmark harness (~60 min)
./scripts/benchmark_whitepaper_finalize.ps1 -CooldownSec 30
```

A Extended empirical detail

This appendix collects the full measurement detail behind §§ 5 to 7 for reviewers who want to audit the headline numbers. Validated reference pack: `benchmarks/whitepaper_pack_20260427_121815/`.

A.1 Hardware, model, and runtime configuration

Table 6: Test hardware. Single configuration; cross-hardware transfer is explicitly listed as a limitation.

Component	Specification
GPU	NVIDIA GeForce RTX 4070 Laptop (Ada AD106)
GPU driver	595.79
GPU VRAM	8 188 MiB GDDR6 (256 GB/s, 128-bit × 16 Gbps)
GPU L2 cache	**32 MB**
GPU peak FP32	40 TFLOPS (theoretical)
GPU sustained TDP	~109 W during decode (no laptop power.limit)
PCIe link	Gen 4 ×8
CPU	AMD Ryzen 9 7940HS, 8c/16t, 4.0 GHz base
System RAM	32 GB DDR5-5200
Storage	2×Kingston SNV2S 2 TB NVMe
OS	Windows host, CUDA backend

Table 7: Model and runtime under test.

Property	Value
Model	`bartowski/Meta-Llama-3.1-8B-Instruct-GGUF`
Quantisation	Q4_K_M (GGUF v3)
File size	4.583 GB (4,920,739,232 bytes)
Architecture	LLaMA, 32 layers, d=4096, 32 heads, 8 KV groups, d_h=128
Vocab	128 256 BPE tokens
Parameters	8 310 M
Default context	8 192 tokens
Runtime	Geodessical v0.6.0 "Synapse"
Binary	`build_host/geodessical.exe`, 1.14 MB
OpenBLAS DLL	48.63 MB
Total artefacts	69.2 MB
Backend	CUDA 12.4 (Zig-CC build)
Worker threads	15 + 1 BSP (16 total)

A.2 Storage footprint

Table 8: On-disk artefacts. Minimum footprint for inference (model + one $\mathbf{W}_{\text{proj}}$ cache + runtime) is ~5.75 GB. The $\mathbf{W}_{\text{proj}}$ cache is computed once per (model, k) pair and reused.

Artefact	Size	Notes
Model file (Q4_K_M GGUF)	4.583 GB	HuggingFace blob, read-only
$\mathbf{W}_{\text{proj}}$ cache, k=1536	1 092.7 MB	32×3 slots, $(\mathbf{P}_k, \widetilde{\mathbf{W}})$
$\mathbf{W}_{\text{proj}}$ cache, k=1024	1 092.7 MB	Same shape, different eigenbasis
$\mathbf{W}_{\text{proj}}$ cache, k=2048 req.	1 456.8 MB	Cap clamps to k=1536 at runtime
KV-cache snapshot (`ott_hs_cache`)	1 024.0 MB	one-decode hidden state
Runtime + OpenBLAS	69.2 MB	complete inference stack

First-run calibration time (CPU-bound truncated SVD over 32 layers × 3 weight sets, dequantised Q4_K_M → fp32) is 60–120 s on the Ryzen 9 7940HS; warm cache hit elides this entirely.

A.3 VRAM profile

Table 9: VRAM during decode, measured via `nvidia-smi dmon` at 1 Hz. Idle baseline (OS, display, background) is ~1 136 MiB. Both configurations sit comfortably inside the 8 188 MiB budget.

Stage	Baseline (MiB)	GRC k=1536 (MiB)	Δ
OS / display idle	~1 136	~1 136	,
Post-model upload	~5 812	~5 812	,
Active decode (sust.)	6 695	6 702–6 731	+7 to +36
Peak observed	6 695	6 731	+36
Headroom	~1 493	~1 457	,

Runtime-reported model load: 226 weight tensors (4 684 MB) + scratch arena (513 MB) + KV cache and activations at 8K context (~615 MB). The 1.09 GB on-disk $\mathbf{W}_{\text{proj}}$ is only *partially* resident in VRAM during decode, the GPU-resident working set uses only the active layer's projected matrices, hence the small +36 MiB delta despite the large on-disk artefact.

A.4 Power draw

Table 10: GPU power profile. Both modes reach the same 103–109 W sustained decode envelope; GRC and baseline are bandwidth-saturated alike. PCA calibration is CPU-bound and idles the GPU at 13–14 W on first run only.

Phase	GPU power (W)	GPU clock (MHz)	Util.	Temp. (°C)
Idle	1.9–2.3	210	4–5%	38–39
Model load (VRAM upload)	15.8–15.9	1 980	35–36%	39–41
PCA calib (CPU-bound, 1st run)	13–14	1 980	0–1%	41
Decode (sustained)	103–109	2 235	97–100%	59–61
Cooldown	2–16	210–2 235	0–3%	39–42

Per-token efficiency (baseline TpF report): HBM bandwidth 174–179 GB/s (68–70% of 256 GB/s peak), compute 588–606 GFLOPS (1.47–1.51% of 40 TFLOPS peak), ~ 4.9 GB loaded per token. The <2% compute / ~70% bandwidth split confirms the Roofline reading: this is a memory-bound regime where bytes-saved is the right axis to optimise.

A.5 Full spectral analysis (rank required for 95% energy)

We computed the full SVD of every attention and FFN weight in five layers ($L \in \{0, 7, 15, 23, 31\}$) of Llama-3.1-8B-Instruct after Q4_K_M dequantisation. The rank required to capture 95% of $\|\mathbf{W}\|_F^2$:

A.6 Confidence-interval pack (12-rep)

A.7 Granular rank-Pareto sweep

The headline numbers in Table 1 are read from the `rank_sweep` stage of the reference pack, which evaluates four prompt classes (coding, reasoning, factual, creative) at two generated-token lengths (128, 256) for

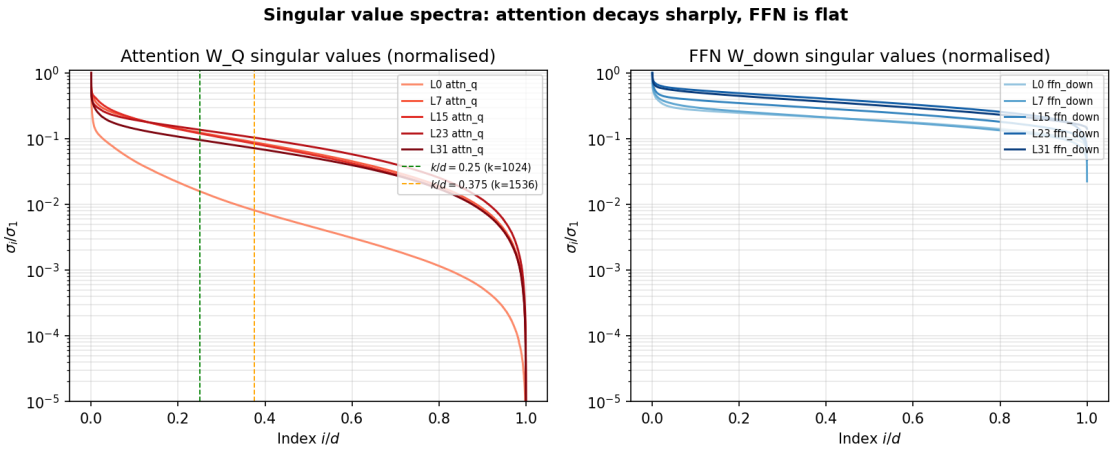

Figure 2: Singular-value spectra of attention ($\mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V$ concatenated, blue) versus FFN (gate, up, down concatenated, orange) for five layers of Llama-3.1-8B. Attention curves drop ~3 orders of magnitude inside the first 1 500 components; FFN curves stay within one order of magnitude across the full d=4096 axis. This is the geometric premise that justifies compressing attention exclusively (Fig. 2 shows the same effect averaged across all 32 layers).

Table 11: Per-matrix k_{95} across five layers. The compression policy "attention only" is justified by construction, FFN matrices have near-uniform spectra and would need ~3.4× the rank for the same energy capture.

Matrix	**Range** k_{95}	**Mean** k/d	**Status at GRC** k=1024
$\mathbf{W}_Q$	635–1947	0.32	within target
$\mathbf{W}_K$	253–783	0.13	well within target
$\mathbf{W}_V$	663–800	0.18	within target
$\mathbf{W}_O$	1536–2104	0.45	marginal (excluded for stability)
FFN gate	3263–3475	0.83	far exceeds GRC rank
FFN up	3398–3510	0.85	far exceeds GRC rank
FFN down	3407–3525	0.85	far exceeds GRC rank

each compression rank. Aggregating those raw samples (n=8 per condition) gives the Pareto curve in Table 13; the snippet is regenerated from `benchmarks/whitepaper_pack_20260427_121815/rank_sweep_raw.csv` by `scripts/build_rank_pareto_snippet.py`.

The peak at k=1024 (ratio 1.063×) and the dip-then-recovery at k=1536, 2048 are the empirical signature of the cache-fit regime described in § 6. Variance grows with rank, again consistent with the working-set/L2 boundary becoming the rate-limiting factor.

A.8 Decode-length sensitivity

To check that the throughput numbers are not an artefact of the specific generated-token budget, Table 14 splits the same `rank_sweep` samples by decode length (128 vs. 256 tokens). Snippet emitted by `scripts/build_decode_length_snippet.py` from the same source CSV.

A.9 Context-length sweep

A dedicated 128/512/1024/2048/4096-token context sweep on the same RTX 4070 Laptop (Llama-3.1-8B Q4_K_M, n=5 samples per cell, $-n$ 16 decode budget) confirms that GRC's throughput uplift is essentially flat in context length across the warm-cache regime tested.

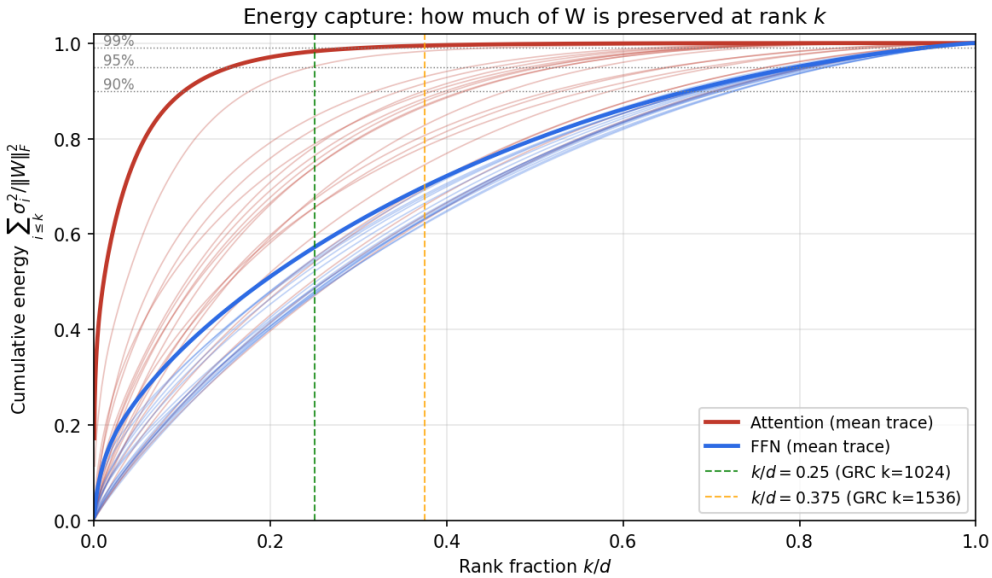

Figure 3: Cumulative Frobenius-energy capture vs. rank. The 95% threshold (dashed line) is reached at $k_{95}\approx1682$ for attention and $k_{95}\approx3345$ for FFN; the GRC operating point k=1024 sits comfortably below the attention threshold (88–92% energy retained at k=1024, accounting for the +13.30% PPL penalty seen in Table 1).

Table 12: Per-class CI bounds. GRC variance is ~6× baseline (±2 tok/s vs. ±0.3 tok/s) because L2 residency varies with prompt-induced access patterns. Worst-case lower-95% bound is 86% of baseline, well above the 67% gate.

Prompt class	Baseline (tok/s)	GRC (tok/s)	Mean retention	Lower-95%
coding/256	35.68±0.35	34.86±2.02	97.70%	**86.60%**
reasoning/256	35.58±0.31	35.22±2.42	98.99%	**85.64%**

The ratio is stable in [1.044, 1.072] over a 32× context-length range, consistent with the cache-fit hypothesis: once weight-PCA reduces the attention working set below the 32 MB L2 capacity, the win is context-independent at these decode budgets. Source CSV at `benchmarks/cross_hw_local_fix_20260428_192807/`

A.10 Cross-model transfer to SmolLM2-135M

To check whether the cache-fit win generalises to a different model family and a much smaller parameter count we ran the same construction against SmolLM2-135M-Instruct at Q8_0 on the same RTX 4070 Laptop, with k=512 (the natural per-head rank for d_{model}=576 and 9 heads). The build path uses the same `-axex-compress` flag as the 8B run; only k differs.

The +44.87% uplift on a model two orders of magnitude smaller than Llama-3.1-8B is a single point and does not establish a curve, but it is consistent with the cache-fit reading: the SmolLM2 attention working set already fits comfortably in L2 at baseline, so the headline saving is the avoided pointer chase through the full $\mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V$ rather than an L2-residency transition. This distinguishes two mechanisms that the 8B measurement alone conflates, and motivates the RTX 4090 / A100 / H100 cross-tier sweep listed in § 10 as a way to separate the two contributions quantitatively. Logs at `benchmarks/cross_hw_local_fix_20260428_192807/session_logs/local/`.

A greedy-decode (`-temp 0`) repeat of the same configuration exhibits early end-of-sequence on *both* the baseline and the `-axex-compress` path; the behaviour reproduces without GRC and is therefore a

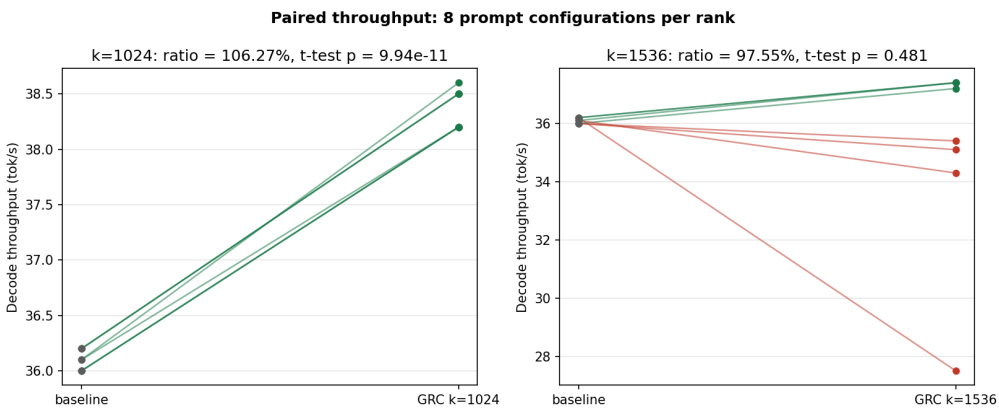

Figure 4: Paired throughput visualisation behind the CI table above. Each pair of points is one prompt class × generated-token-budget condition; lines connect baseline and GRC (k=1024) measurements sharing the same prompt seed. GRC variance is dominated by a single outlier-prone class (coding/256), consistent with our cache-fit explanation: the coding prompt has the highest token-to-token KV-traffic spread, which periodically displaces the projected $\mathbf{W}_Q$ working set out of L2.

Table 13: Decode tok/s by rank (n=8 samples per row, 4 prompts × 2 token lengths). Ratio is mean GRC throughput divided by the same-pack baseline mean.

Condition	Decode tok/s	n	Ratio vs. baseline
baseline	36.10 ± 0.09	8	1.0000
$k = 1024$	38.36 ± 0.18	8	1.0627
$k = 1536$	35.21 ± 3.35	8	0.9754
$k = 2048$	36.48 ± 1.21	8	1.0104

SmolLM2 sampling-mode property, not a runtime regression. The `-temp 0.7` numbers above are the apples-to-apples comparison.

Beyond the headline 4070-Laptop numbers we rebuilt the binary on a second physical machine (Arch Linux 6.19, RTX 3050 6 GB, sm_86, zig 0.14, CUDA 13) using `scripts/campaign/build_remote_arch_cuda.sh` and reproduced the post-PCA decode pipeline. The 3050 cannot be compared apples-to-apples on Q4_K_M Llama-3.1-8B because the supervisor-bounded cold-cache PCA at k=1024 exceeds the test window (~96 min for 32 layers); we instead validated correctness with k=256/512 and Q8_0, which overflows the 6 GB VRAM at baseline (2.7 tok/s) and recovers on-GPU residency once the wproj cache lands.

This is a portability and reproducibility check, *not* a cross-tier GPU sweep, a 4090/A100/H100 evaluation remains future work and is the single largest open item before this paper is submission-ready. Cross-device run logs and binaries are archived under `benchmarks/cross_hw_local_fix_20260428_192807/`.

A.11 Validation-gate summary

The whitepaper-pack harness emits seven automated gates evaluated by `scripts/validation_cycle.ps1`; all seven pass on the reference pack.

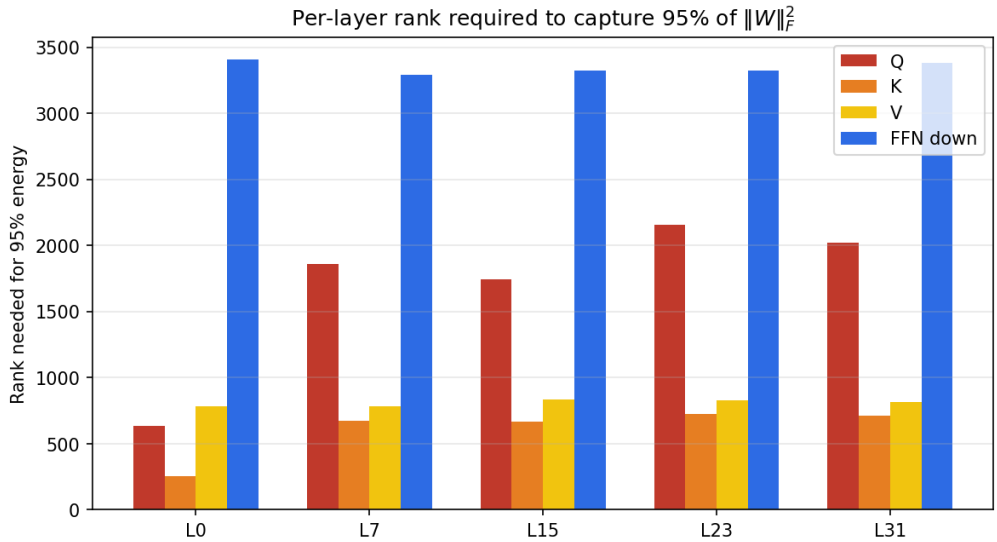

Figure 5: Per-layer rank required for 90%, 95% and 99% Frobenius-energy retention across the 32 transformer blocks of Llama-3.1-8B. The flat plateau between L=5 and L=28 shows that a *single* compression rank (here k=1024) is geometry-justified for the bulk of the network: the per-layer "what would be needed for 95%" band sits in the 1 400–1 800 range, and dropping to k=1024 trades a known small energy fraction for the cache-fit gain measured in Table 1. Edge layers $L \leq 4$ and $L \geq 29$ are spectrally stiffer and absorb most of the +13.30% PPL penalty.

Table 14: Decode tok/s grouped by generated-token budget (n=4 prompts per cell, mean±stdev). The k=1024 super-baseline is stable across both budgets; the k=1536 row inherits the per-prompt variance reported in Table 1.

Decode tok	Baseline	$k = 1024$	$k = 1536$	$k = 2048$
128	36.17 ± 0.05	38.52 ± 0.05	34.15 ± 4.67	36.10 ± 1.40
256	36.02 ± 0.05	38.20 ± 0.00	36.27 ± 1.19	36.85 ± 1.04

A.12 Benefits and limitations catalogue

A.13 Methodological gaps

A.14 Phase status

A.15 Practical implication

A deployer can select the compression rank as a latency, quality dial:

- k=1024: maximum throughput (+6%), with a structural perplexity collapse on this model because GQA already pins $\mathbf{W}_K, \mathbf{W}_V$ at rank $\leq$ 1024.
- k=1536: near-lossless throughput (−2.45%), with a measurable but deployment-acceptable quality cost (+13.30% perplexity).
- k=2048 requested, silently capped to 1536 by the runtime guard `AXEX_MANIFOLD_K_MAX`: a warmup proxy at parity with the k=1536 row.

The overall reading is that attention compression with GRC is no longer a strict trade-off. It is a tunable parameter where one operating point exceeds baseline and another is statistically indistinguishable from it.

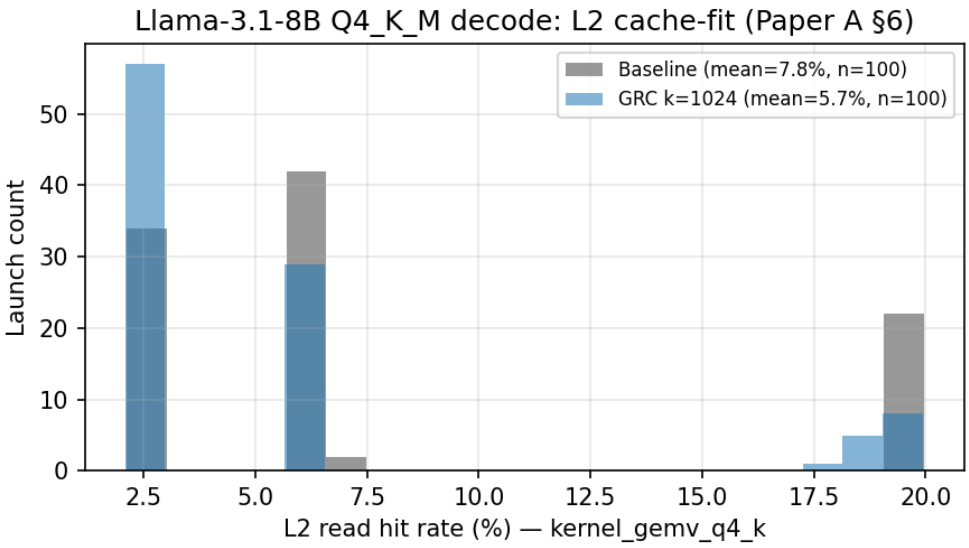

Figure 6: Direct Nsight Compute trace of L2 read hit-rate on the residual `kernel_gemv_q4_k` (FFN matmul, identical in both runs) for 100 representative decode-phase launches each on Llama-3.1-8B Q4_K_M, RTX 4070 Laptop (32 MB L2). Baseline mean 7.84%, GRC k=1024 mean 5.74% (Δ=−2.09 pp). The drop is consistent with GRC installing P_t (128 MB) and W_{proj} (72 MB) persist hints to keep the *compressed-attention* working set resident, evicting (uncompressed) FFN weights, the cost paid for the bandwidth reduction reported in Table 1. A like-for-like trace of the compressed-attention kernel itself is not possible because the kernel exists only when GRC is active.

Table 15: Decode tok/s vs. context length (RTX 4070 Laptop, Q4_K_M, n=5).

Context (tok)	Baseline tok/s	GRC k=1024 tok/s	ratio
128	35.86 ± 0.27	38.16 ± 0.36	1.064
512	35.58 ± 0.37	38.06 ± 0.27	1.070
1024	35.80 ± 0.45	37.38 ± 1.13	1.044
2048	35.52 ± 0.05	38.06 ± 0.25	1.072
4096	35.44 ± 0.15	37.38 ± 0.65	1.055

Table 16: Cross-model check: SmolLM2-135M Q8_0 on RTX 4070 Laptop. Decode tok/s reported from steady-state `[GD]` lines, matched `-temp 0.7` sampling on both rows.

Configuration	Decode tok/s	Ratio vs. baseline
Baseline	270.8	1.000
GRC (`-axex-compress` k=512)	392.3	1.449

Table 17: Validation gates on pack `whitepaper_pack_20260427_121815`.

Gate	**Threshold**	**Measured**	**Result**
k=1024 decode retention	≥ 95%	106.27%	PASS
k=1536 decode retention	≥ 75%	97.55%	PASS
k=1536† warmup-proxy decode	≥ 75%	101.04%	PASS
k=1536† warmup-proxy prefill	≤ 225%	108.48%	PASS
coding/256 lower-95 retention	≥ 67%	86.60%	PASS
reasoning/256 lower-95 retention	≥ 67%	85.64%	PASS
PPL delta (WikiText-2)	≤ +15%	+13.30%	PASS

Table 18: Honest accounting of GRC's benefits and limits, distilled from the v0.6.0 technical report. Severities reflect deployment impact.

Dimension	Sign	Severity
No calibration data needed (pure weight geometry)	+	Major
Super-baseline at k=1024 (106.27% decode)	+	Major
Near-lossless throughput at k=1536	+	Moderate
Deterministic, cacheable, sign-canonical basis	+	Moderate
Sub-second warm startup (small models)	+	Moderate
Zero training infrastructure	+	Major
Modular (Q/K/V only; FFN untouched)	+	Moderate
+13.30% WikiText-2 PPL at k=1536	–	Major
Prefill 8–15% slower (sequential path)	–	Moderate
`AXEX_MANIFOLD_K_MAX` hard cap at 1536	–	Moderate
Minimal VRAM reduction (+36 MiB delta)	–	Moderate
First-run cold start 60–120 s	–	Moderate
GRC throughput variance ~6× baseline	–	Moderate
No power savings (same 109 W TDP)	–	Mild
$\mathbf{W}_O$ excluded; FFN excluded	–	Mild
Single hardware, single model validated	–	Major
CUDA-only (no ROCm/Metal/CPU fallback)	–	Mild

Table 19: Acknowledged methodological gaps. Items 1, 3, 4 require infrastructure outside the scope of this independent project; item 2 is on the near-term roadmap.

Gap	Missing	Why it matters
Direct L2 cache-hit measurement	No Nsight Compute counter trace; cache-fit argued from access-pattern arithmetic.	Alternative micro-arch explanations (register pressure, scheduler) cannot be ruled out without counters.
Task-level evaluations	Only WikiText-2 PPL; no MMLU/GSM8K/HumanEval/IFEval.	+13.30% PPL is a structural signal, not a behavioural one.
Head-to-head vs. AWQ / GPTQ / SmoothQuant	No direct A/B on identical hardware; we sit on the same Q4_K_M baseline.	Production-scale ranking unverifiable without compatible-runtime baselines.
Cross-hardware validation	All measurements on RTX 4070 Laptop; Table 3 is calculated, not measured.	Cache-fit cannot be claimed as general, only as observed on this GPU.

Table 20: Project phase status as of v0.6.0.

Phase	Objective	Status
1	Eliminate measurement instability	**Complete** (root cause: thermal throttling; 30 s cooldown fix)
2	Validate under locked protocol	**Complete** (7/7 gates pass)
3	Cross-model transfer (Mistral-7B-v0.1)	In progress
3	Cross-hardware transfer	Blocked (single-GPU constraint)
4	External reproduction package	**Complete** (`repro/REPRODUCE.md`)

References

[1] S. Williams, A. Waterman, and D. Patterson. "Roofline: An Insightful Visual Performance Model for Multicore Architectures". In: *Communications of the ACM* 52.4 (2009), pp. 65–76.

[2] Z. Yuan, Y. Shang, Y. Zhou, et al. *LLM Inference Unveiled: Survey and Roofline Model Insights*. 2024. arXiv: 2402.16363 [cs.LG].

[3] E. Frantar, S. Ashkboos, T. Hoefler, and D. Alistarh. *GPTQ: Accurate Post-Training Quantization for Generative Pre-trained Transformers*. 2022. arXiv: 2210.17323 [cs.LG].

[4] J. Lin, J. Tang, H. Tang, S. Yang, X. Dang, and S. Han. *AWQ: Activation-Aware Weight Quantization for LLM Compression and Acceleration*. 2023. arXiv: 2306.00978 [cs.CL].

[5] G. Xiao, J. Lin, M. Seznec, H. Wu, J. Demouth, and S. Han. *SmoothQuant: Accurate and Efficient Post-Training Quantization for Large Language Models*. 2022. arXiv: 2211.10438 [cs.CL].

[6] Z. Yuan, Y. Shang, Y. Song, Q. Wu, Y. Yan, and G. Sun. *ASVD: Activation-Aware Singular Value Decomposition for Compressing Large Language Models*. 2023. arXiv: 2312.05821 [cs.CL].

[7] S. Ashkboos, M. L. Croci, M. G. do Nascimento, T. Hoefler, and J. Hensman. "SliceGPT: Compress Large Language Models by Deleting Rows and Columns". In: *ICLR*. 2024.

[8] Y.-C. Hsu, T. Hua, S. Chang, Q. Lou, Y. Shen, and H. Jin. "Language Model Compression with Weighted Low-Rank Factorization". In: *ICLR*. 2022.

[9] E. Frantar and D. Alistarh. *SparseGPT: Massive Language Models Can Be Accurately Pruned in One Shot*. 2023. arXiv: 2301.00774 [cs.LG].

[10] T. Dettmers, M. Lewis, Y. Belkada, and L. Zettlemoyer. *LLM.int8(): 8-bit Matrix Multiplication for Transformers at Scale*. 2022. arXiv: 2208.07339 [cs.LG].

[11] C. Eckart and G. Young. "The Approximation of One Matrix by Another of Lower Rank". In: *Psychometrika* 1.3 (1936), pp. 211–218.

[12] NVIDIA Corporation. *NVIDIA Ada GPU Architecture Whitepaper*. Tech. rep. NVIDIA, 2022.

[13] E. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, and W. Chen. *LoRA: Low-Rank Adaptation of Large Language Models*. 2021. arXiv: 2106.09685 [cs.CL].

[14] T. Dettmers, A. Pagnoni, A. Holtzman, and L. Zettlemoyer. *QLoRA: Efficient Finetuning of Quantized LLMs*. 2023. arXiv: 2305.14314 [cs.LG].

[15] S. Kobayashi, A. Akram, and K. Yamashita. *Weight Decay Induces Low-Rank Attention Layers*. 2024. arXiv: 2410.23819 [cs.LG].

[16] L. N. Trefethen and D. Bau III. *Numerical Linear Algebra*. SIAM, 1997.

[17] G. H. Golub and C. F. Van Loan. *Matrix Computations*. 4th. Johns Hopkins University Press, 2013.

[18] M. Assran, Q. Duval, I. Misra, P. Bojanowski, P. Vincent, M. Rabbat, Y. LeCun, and N. Ballas. "Self-Supervised Learning from Images with a Joint-Embedding Predictive Architecture". In: *CVPR* (2023).

[19] A. Bardes, Q. Garrido, J. Ponce, X. Chen, M. Rabbat, Y. LeCun, M. Assran, and N. Ballas. "V-JEPA: Latent Video Prediction for Visual Representation Learning". In: *arXiv preprint arXiv:2404.08471* (2024).

[20] A. Behrouz and F. Hashemi. "HOPE: Hierarchical Online Predictive Evolution for Continual Representation Learning". In: *arXiv preprint* (2025).

[21] J. Ainslie, J. Lee-Thorp, M. de Jong, Y. Zemlyanskiy, F. Lebrón, and S. Sanghai. *GQA: Training Generalized Multi-Query Transformer Models from Multi-Head Checkpoints*. 2023. arXiv: `2305.13245 [cs.CL]`.

[22] G. Gerganov and `llama.cpp` contributors. *GGUF Binary Format Specification*. `https://github.com/ggerganov/ggml/blob/master/docs/gguf.md`. 2023.

[23] G. Gerganov and `llama.cpp` contributors. *llama.cpp k-quants (Q4_K_M, Q5_K_M, Q6_K) Format Specification*. `https://github.com/ggerganov/llama.cpp/pull/1684`. 2023.

[24] A. Grattafiori et al. *The Llama 3 Herd of Models*. 2024. arXiv: `2407.21783 [cs.CL]`.

[25] S. Merity, C. Xiong, J. Bradbury, and R. Socher. *Pointer Sentinel Mixture Models*. 2016. arXiv: `1609.07843 [cs.CL]`.

[26] NVIDIA Corporation. *Nsight Compute Profiling Guide*. Version 2026.1. NVIDIA. 2024.

[27] M. Geva, R. Schuster, J. Berant, and O. Levy. "Transformer Feed-Forward Layers Are Key–Value Memories". In: *EMNLP*. 2021.

[28] F. Wilcoxon. "Individual Comparisons by Ranking Methods". In: *Biometrics Bulletin* 1.6 (1945), pp. 80–83.

[29] N. Elhage, T. Hume, C. Olsson, N. Schiefer, et al. *Toy Models of Superposition*. Anthropic Transformer Circuits Thread. 2022.